

Classical statistical inference

Iterated expectations¹

- a. Given $E[X|Y] = 2Y$ and $f(Y) = .5$ with $Y \in [-3, -1]$, what is $E[X]$?

Solution:

$$\begin{aligned} E[X] &= E_Y[E_{Y|X}[X|Y]] \\ &= \int_{-3}^{-1} E[X|Y] \Pr(Y) dy \\ &= \int_{-3}^{-1} 2 \times Y (.5) dy \\ &= \frac{Y^2}{2} \Big|_{-3}^{-1} \\ &= \frac{1}{2} - \frac{9}{2} \\ &= -\frac{8}{2} \\ &= -4 \end{aligned}$$

- b. Given $E[Z|H] = 15H - 10$ and $H \sim \text{Bernoulli}(.2)$, what is $E[Z]$?

Solution:

$$\begin{aligned} E[Z] &= E_H[E_{Z|H}[Z|H]] \\ &= \sum_h E[Z|H = h] p(h) \\ &= E[Z|H = 1] \Pr(H = 1) + E[Z|H = 0] \Pr(H = 0) \\ &= 5 \times .2 + -10 \times .8 \\ &= 1 - 8 \\ &= -7 \end{aligned}$$

Properties of estimators

- a. Let $X_1, \dots, X_n \sim \text{Poisson}(\lambda)$ and let $\hat{\lambda} = \frac{\sum_{i=1}^n X_i}{n}$. Find the bias, standard error, and MSE of this estimator.²

- b. Bias

Recall that $\text{bias}(\hat{\lambda}) = E_{\lambda}[\hat{\lambda}] - \lambda$. The expected value of the estimator $\hat{\lambda} = E[X] = \lambda$. Because the expected value of the estimator $\hat{\lambda}$ is the true value λ , $\text{bias} = 0$.

- a. Standard error

¹Grimmer 13.1

²Wasserman 6.6.1

$$\text{Var}(\hat{\lambda}) = \frac{\text{Var}[X]}{n} = \frac{\lambda}{n}$$

$$se = \sqrt{\frac{\lambda}{n}}$$

a. MSE

$$\begin{aligned}\text{MSE} &= \text{bias}^2(\hat{\lambda}) + \text{Var}_{\lambda}(\hat{\lambda}) \\ &= 0^2 + \frac{\lambda}{n} \\ &= \frac{\lambda}{n}\end{aligned}$$

a. Let $X_1, \dots, X_n \sim \text{Uniform}(0, \theta)$ and let $\hat{\theta} = 2\bar{X}_n$. Find the bias, standard error, and MSE of this estimator.³

Recall that for the uniform distribution

$$E(X_i) = \frac{\theta}{2}, \text{Var}(X_i) = \frac{\theta^2}{12}$$

a. Bias

$$E_{\theta}(2\bar{X}) = 2E_{\theta}(\bar{X}) = 2 \times \frac{\theta}{2} = \theta$$

$$\text{bias}(\hat{\theta}) = E_{\theta}[\hat{\theta}] - \theta = \theta - \theta = 0$$

a. Standard error

$$\text{Var}_{\theta}(2\bar{X}) = 4\text{Var}_{\theta}(\bar{X}) = \frac{4\sigma^2}{12n} = \frac{4\sigma^2}{12n} = \frac{\sigma^2}{3n}$$

$$se = \sqrt{\text{Var}_{\theta}(2\bar{X})} = \sqrt{\frac{\sigma^2}{3n}}$$

a. MSE

$$\begin{aligned}\text{MSE} &= \text{bias}^2(\hat{\theta}) + \text{Var}_{\theta}(\hat{\theta}) \\ &= 0^2 + \frac{\sigma^2}{3n} \\ &= \frac{\sigma^2}{3n}\end{aligned}$$

³Wasserman 6.6.3

Birds of a feather get their news on Twitter (X)

A poll conducted in 2013 found that 52% of U.S. adult Twitter users get at least some news on Twitter. The standard error for this estimate was 2.4%, and a normal distribution may be used to model the sample proportion.⁴

- a. Construct a 99% confidence interval for the fraction of U.S. adult Twitter users who get some news on Twitter, and interpret the confidence interval in context.

Solution: Recall the general formula is Point estimate $\pm z^* \times SE$. First, identify the three different values. The point estimate is 52%, $z^* = 2.58$ for a 99% confidence level, and $SE = 2.4\%$. Then plug the values into the formula:

$$52\% \pm 2.58 \times 2.4\% \rightarrow (45.8\%, 58.4\%)$$

(T/F) Birds of a feather get their news on Twitter (X)

Identify the following statements as true or false using the above information. Provide an explanation to justify each of your answers.

- a. The data provide statistically significant evidence that more than half of U.S. adult Twitter users get some news through Twitter. Use a significance level of $\alpha = 0.01$.

Solution: False. 50% is included in the 99% confidence interval, hence a null hypothesis of $p = 0.50$ would not be rejected at this level.

- a. Since the standard error is 2.4%, we can conclude that 97.6% of all U.S. adult Twitter users were included in the study.

Solution: False. The standard error measures the variability of the sample proportion, and is unrelated to the proportion of the population included in the study. \rightarrow

- a. If we want to reduce the standard error of the estimate, we should collect less data.

Solution: False. We need to increase the sample size to decrease the standard error.

- a. If we construct a 90% confidence interval for the percentage of U.S. adults Twitter users who get some news through Twitter, this confidence interval will be wider than a corresponding 99% confidence interval.

Solution: False. As the confidence level decreases so does the margin of error, and hence the width of the confidence interval.

Choose your own death

There is a theory that people can postpone their death until after an important event. To test the theory, Phillips and King (1988) collected data on deaths around the Jewish holiday Passover. Of 1919 deaths, 922 died the week before the holiday and 997 died the week after. Think of this as a binomial and test the null hypothesis that $\theta = \frac{1}{2}$. Report and interpret the p -value. Also construct a confidence interval for θ .⁵

Solution: $\hat{\theta} = \frac{922}{1919} = 0.48$ with estimated standard error $\hat{se} = \sqrt{\frac{\hat{p}(1-\hat{p})}{n}} = 0.011$. The Wald test statistic is

$$\frac{0.48 - 0.50}{0.011} = -1.818$$

⁴OI 4.8 and 4.10

⁵Wasserman 10.11.6

and the p -value is $\Pr(|Z| > 1.818) = 2\Pr(Z > 1.818) = 0.07$ which is moderate evidence against the null $H_0 : p = \frac{1}{2}$. A 95% confidence interval is

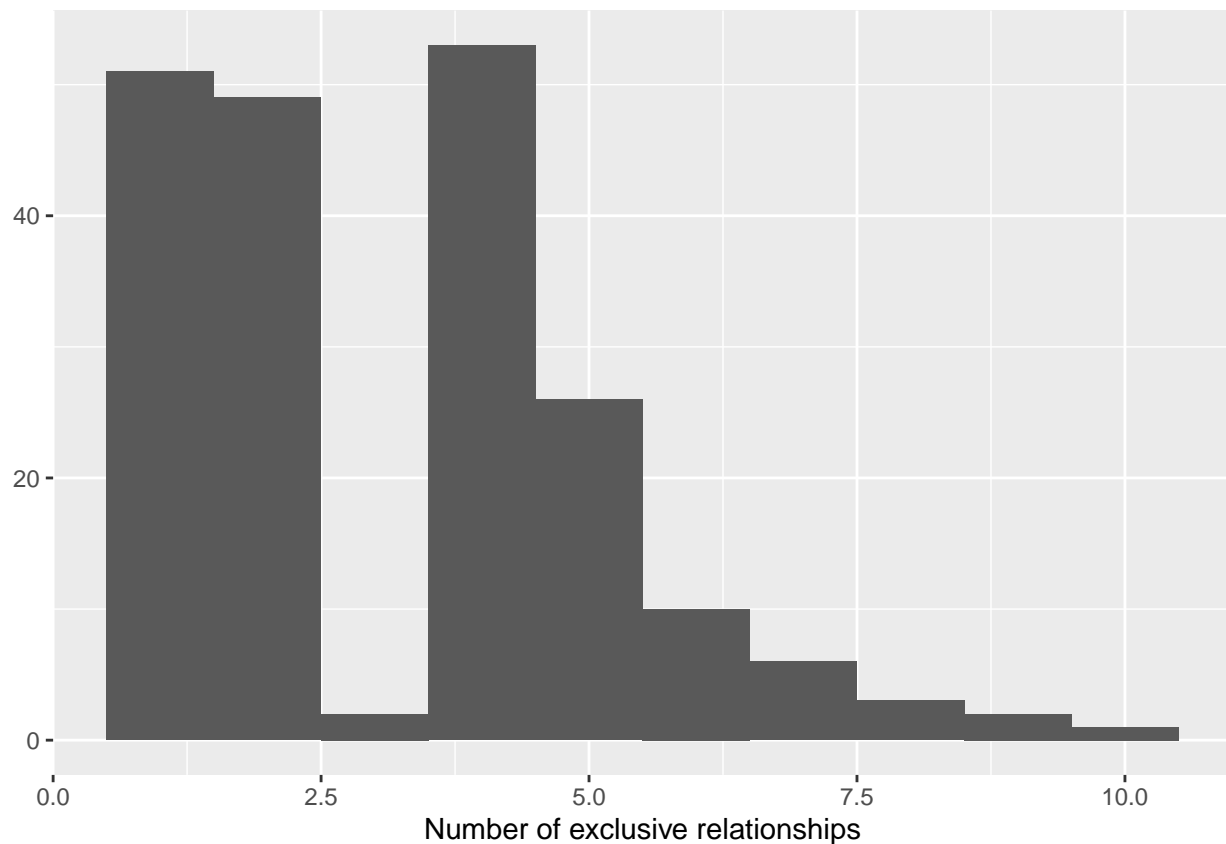
$$\hat{p} \pm 2 \times \hat{se} = (0.458, 0.503)$$

The results are rather equivocal. I would suggest collecting further data.

Dating on college campuses

A survey conducted on a reasonably random sample of 203 undergraduates asked, among many other questions, about the number of exclusive relationships these students have been in. The histogram below shows the distribution of the data from this sample.

The sample average is 3.2 with a standard deviation of 1.97.



Estimate the average number of exclusive relationships undergraduate students have been in using the Normal distribution and a 90% confidence interval and interpret this interval in context.⁶

Solution: The confidence interval can be calculated as follows:

$$\begin{aligned} \bar{x} \pm Z^* SE_{\bar{x}} &= 3.2 \pm 1.65 \times \frac{1.97}{\sqrt{203}} \\ &= 3.2 \pm 0.23 \\ &= (2.97, 3.43) \end{aligned}$$

⁶OI 4.15

Statistical significance

Determine whether the following statement is true or false, and explain your reasoning: “With large sample sizes, even small differences between the null value and the point estimate can be statistically significant.”⁷

Solution: True. If the sample size is large, then the standard error will be small, meaning even relatively small differences between the null value and point estimate can be statistically significant. **This is kind of a big deal when analyzing big data.**

Sleep deprivation

New York is known as “the city that never sleeps”. A random sample of 25 New Yorkers were asked how much sleep they get per night. Statistical summaries of these data are shown below. Do these data provide strong evidence that New Yorkers sleep less than 8 hours a night on average?⁸

n	\bar{x}	s	min	max
25	7.73	0.77	6.17	9.78

- a. Write the hypotheses in symbols and in words.

Solution:

$H_0 : \mu = 8$ (New Yorkers sleep 8 hrs per night on average.)

$H_A : \mu \neq 8$ (New Yorkers don’t sleep 8 hrs per night on average.)

- b. Calculate the test statistic, T , and the associated degrees of freedom.

Solution: The test statistic and degrees of freedom can be calculated as follows:

$$T = \frac{\bar{x} - \mu_0}{\frac{s}{\sqrt{n}}} = \frac{7.73 - 8}{\frac{0.77}{\sqrt{25}}} = \frac{-0.27}{0.154} = -1.75$$

and

$$df = 25 - 1 = 24$$

- a. Find and interpret the p-value in this context.

Solution: p-value = $\Pr(T_{24} < -1.75) = 0.095$. If in fact the true population mean of the amount New Yorkers sleep per night was 8 hours, the probability of getting a random sample of 25 New Yorkers where the average amount of sleep is farther from the mean than 7.73 hrs per night is about 0.09.

- a. What is the conclusion of the hypothesis test?

Solution: Since p-value > 0.05 , FAIL to reject H_0 . The data DO NOT provided convincing evidence that New Yorkers sleep less than 8 hours per night on average.

- a. If you were to construct a 90% confidence interval that corresponded to this hypothesis test, would you expect 8 hours to be in the interval?

Solution: No, the hypothesis test suggests that the average amount of sleep New Yorkers get is lower than 8 hours per night, therefore we wouldn’t expect 8 hours to be in the interval. Note that if we were using a one-tailed test with $\alpha = 0.05$, we would incorrectly reject the null.

⁷OI 4.47

⁸OI 5.7

Interpreting public opinion polls

On June 28, 2012 the U.S. Supreme Court upheld the much debated 2010 healthcare law, declaring it constitutional. A Gallup poll released the day after this decision indicates that 46% of 1,012 Americans agree with this decision. At a 95% confidence level, this sample has a 3% margin of error. Based on this information, determine if the following statements are true or false, and explain your reasoning.⁹

- a. We are 95% confident that between 43% and 49% of Americans in this sample support the decision of the U.S. Supreme Court on the 2010 healthcare law.

Solution: False. A confidence interval is constructed to estimate the population proportion, not the sample proportion.

- a. We are 95% confident that between 43% and 49% of Americans support the decision of the U.S. Supreme Court on the 2010 healthcare law.

Solution: True. This is the correct interpretation of the confidence interval, which can be calculated as $0.46 \pm 0.03 = (0.43, 0.49)$. A nicer restatement might be 'we are 95% confident that the proportion of Americans who support the decision of the U.S. Supreme Court on the 2010 healthcare law is between 43% and 49%.

- a. If we considered many random samples of 1,012 Americans, and we calculated the sample proportions of those who support the decision of the U.S. Supreme Court, 95% of those sample proportions will be between 43% and 49%.

Solution: Not our best response. The confidence interval does not tell us what we might expect to see in another random sample. We COULD say that 95% of our confidence intervals will capture the mean.

- b. The margin of error at a 90% confidence level would be higher than 3%.

Solution: False. As the confidence level decreases, the margin of error decreases as well.

Survey Qs:

1. How prepared do you feel for the final exam? not prepared (1) 2 3 4 (5) very prepared
2. Which content area of the course was the easiest for you? Why? Please be specific.

3. Which content area of the course was the most challenging for you? Why? Please be specific.

⁹OI 6.6